

Portkey AI Builders Challenge

Build Real Production AI Systems

Welcome to the Portkey AI Builders Challenge — a 24-hour hackathon focused on building real, production-grade AI systems, not demos.

This is not about chatbots or flashy UI.

This is about living systems that run continuously, manage state, adapt to change, and make correctness trade-offs.

If your system ran unattended for 6 months:

- would you trust it?
- would an enterprise trust it?

That's the bar.

Format

- Duration: 24 hours
 - Team size: 3 people
 - Teams: ~20
 - Tracks: Open choice (pick any one problem below)
 - Data: Real or synthetic
 - Models & Tools: No restrictions
-

Embrace AI (Important)

We are actively looking for teams that leverage LLMs and AI tools deeply and creatively.

- There is no restriction on:
 - models (OSS or hosted)
 - providers
 - fine-tuning approaches
 - agent frameworks
 - developer AI tools
- You are encouraged to:
 - use LLMs as planners, critics, verifiers, and agents
 - use AI to write code, generate data, validate outputs, and reason about trade-offs
 - automate workflows using AI wherever it makes sense

This challenge is not about “minimising AI usage.”

It’s about using AI well.

🎯 What We’re Looking For

Across all projects, judges will look for:

- Continuous systems (not one-shot scripts)
- Thoughtful use of LLMs and AI tools
- State management & freshness logic

- Clear trade-offs (cost vs quality vs safety)
 - Failure handling
 - Explainability & observability
 - Engineering rigor over UI polish
-



Challenge Tracks

Teams may pick any one of the following problems.

1 Living LLM Pricing & Capability Catalog

Problem

LLM providers constantly change:

- pricing
- context limits
- supported features
- rate limits
- model availability

These changes are fragmented across docs, blogs, APIs, and changelogs — and often not announced clearly.

Goal

Build a living system that continuously tracks LLM providers and models, and maintains a canonical catalog of:

- models
- pricing
- limits
- capabilities
- metadata

Core Requirements

Your system should:

- use agents to ingest data (scraping, APIs, feeds)
- maintain state to avoid unnecessary re-crawls
- detect semantic changes, not just text diffs
- attach metadata to every field:
 - source(s)
 - last verified timestamp
 - confidence score
- surface conflicting information instead of overwriting

Stretch Ideas

- pricing change alerts
 - historical pricing deltas
 - breaking-change detection
 - machine-readable registry (JSON)
-

2 Smart Prompt Parser & Canonicalisation Engine

Problem

AI systems accumulate thousands of prompts:

- duplicated
- slightly modified
- hard to reason about
- impossible to optimise at scale

Goal

Build a system that analyses a collection of prompts and:

1. clusters semantically equivalent prompts
2. extracts prompt templates:
 - constant parts
 - variable slots
3. detects evolution:
 - brand-new prompt families
 - new variable sets in existing families

Requirements

- semantic similarity (not string matching)
- explainable grouping decisions
- incremental processing

- confidence thresholds for merges
-

3 Fine-Tuned OSS Guardrail Models

Problem

Prompt-based guardrails are:

- expensive
- slow
- inconsistent

Goal

Fine-tune one or more open-source models to act as specialised guardrail models.

Use Cases

Teams may fine-tune for one or more of:

- PII detection
- prompt injection detection
- content moderation
- policy violation detection
- unsafe instruction detection

Requirements

- synthetic data generation strategy
- clear train / eval split

- precision-focused evaluation
- runtime inference demo
- discussion of failure cases

Depth > breadth.

4 Cost–Quality Optimisation via Historical Replay

Problem

Model choices are often made blindly.

Goal

Build a system that:

- replays historical prompt–completion data
- evaluates across models and guardrails
- measures cost, quality, refusal rates
- recommends better trade-offs

Expected Output



Judging Criteria

Projects will be evaluated on:

1. Production readiness
2. Thoughtful use of LLMs & AI tools

3. System design & state management
4. Correctness & trade-offs
5. Engineering quality
6. Failure handling
7. Explainability

UI polish is optional.

Clear thinking is not.



Getting Started with Portkey (Recommended)

Before you start building, we strongly recommend that you:

- Go through the Portkey documentation
- Sign up on the Portkey platform
- Play around with the Gateway, models, and configurations

A lot of the things in this challenge — routing across models, guardrails, retries, logging, cost tracking, and experimentation — become significantly easier once you understand how Portkey works.

Additionally, by signing up you will:

- get AI credits across different model providers
 - be able to experiment with multiple models quickly
 - avoid re-building common plumbing from scratch
-



Final Note

This is a builders challenge.

Use AI aggressively.

Automate ruthlessly.

Design systems that survive contact with reality.

Good luck — build something you'd actually ship.